



Please act as professional software architect designer to research on Day 2 research topic, challenge the finds with industrial best practice and provide the best solution after internal debate.

Based on my comprehensive Day 2 research, I've completed an in-depth analysis of agent communication architecture. Here's my professional architectural assessment after challenging findings with industrial best practices:

Agent Communication Architecture: Professional Design Analysis

Executive Summary

After analyzing 45+ authoritative sources covering communication protocols, message bus technologies, schema design, and retry patterns, I recommend **Redis Streams with JSON-based messages** as the primary communication backbone, supplemented by **Redis Pub/Sub for UI updates** and **WebSocket for real-time dashboards**. This hybrid approach balances performance, reliability, and operational simplicity.

Critical Architecture Revision from Day 1

Updated Message Bus Decision: Redis Streams (Not Pub/Sub)

Day 1 Recommendation: Redis Pub/Sub for simplicity and sub-millisecond latency

Day 2 Revision: Redis Streams as primary message bus

Why the Change? Deep-dive research revealed critical production requirements that Pub/Sub cannot satisfy: [\[1\]](#) [\[2\]](#) [\[3\]](#) [\[4\]](#) [\[5\]](#)

Feature comparison between Redis Pub/Sub and Redis Streams for agent messaging architecture

Key Differentiators: [\[1\]](#) [\[3\]](#) [\[5\]](#)

- Message Persistence:** Streams writes to disk; Pub/Sub is ephemeral. For debugging agent coordination failures, message history is non-negotiable.

2. **Consumer Groups:** Streams enables horizontal scaling—multiple instances of the same agent type process messages from one stream without duplication. Pub/Sub broadcasts to all subscribers (no work distribution).^[2]
3. **Exactly-Once Semantics:** Streams with acknowledgment prevents duplicate task execution. Pub/Sub is at-most-once (fire-and-forget).^[3]
4. **Recovery Capability:** Streams allows re-reading unprocessed messages after crashes. Pub/Sub subscribers miss messages if offline.^[4]
5. **Latency Trade-off:** Only 1-5ms vs <1ms—acceptable for test generation workflows (not high-frequency trading).^[6] ^[5]

Performance Reality Check:^[6]

- **Redis Pub/Sub:** Sub-millisecond latency, but zero durability
- **Redis Streams:** 1-5ms latency with full persistence and replay

For a test automation system where **auditability and reliability** trump raw speed, the 1-4ms latency penalty is insignificant compared to the operational benefits.

Communication Protocol Analysis

gRPC vs REST: The Performance Debate

Performance Benchmarks:^[7] ^[8] ^[9] ^[10] ^[11] ^[12]

Performance comparison of communication protocols showing latency (ms) and throughput (operations/sec) for agent messaging

gRPC's 7-10x Speed Advantage Explained:^[9] ^[10]

- **Binary Protocol Buffers** vs JSON text (10x faster serialization)
- **HTTP/2 multiplexing** vs HTTP/1.1 sequential (connection reuse)
- **Predictable latency:** gRPC has no large outliers; REST shows high P99 variability^[9]

Real-World Test Results:^[11] ^[9]

Load Test: 10,000 requests

REST:

- 2,500 req/sec
- Avg latency: 40ms
- P99 latency: 120ms

gRPC:

- 8,000 req/sec (3.2x improvement)
- Avg latency: 12ms
- P99 latency: 35ms (3.4x better)

When NOT to Use gRPC

Despite superior performance, gRPC has critical drawbacks for this use case:[\[10\]](#) [\[11\]](#) [\[12\]](#)

Implementation Overhead:[\[10\]](#)

- **45 minutes** for simple gRPC service vs **minutes** for REST
- Requires Protocol Buffer schema definitions
- Less mature tooling ecosystem
- Steeper learning curve for team

Browser Incompatibility:[\[11\]](#)

- Requires gRPC-Web proxy for browser clients
- WebSocket is native for real-time UI updates

Operational Complexity:[\[12\]](#)

- Debugging binary payloads harder than JSON
- Limited third-party monitoring tools vs REST's Postman, Swagger, etc.

Architecture Decision: Use Message Bus, Not gRPC

Critical Insight: The debate is **not gRPC vs REST** for agent communication—it's **message bus vs direct RPC**.

Why Message Bus Wins for Agent Coordination:

1. **Decoupling:** Agents don't need to know each other's network addresses
2. **Asynchronous:** Non-blocking task delegation (agents pull work when ready)
3. **Buffering:** Queue absorbs bursts (100+ simultaneous requests)
4. **Retry logic:** Message bus handles retries transparently
5. **Observability:** Centralized message tracing without distributed tracing complexity

When to Use Each Protocol:

- **Agent ↔ Agent (Backend):** Redis Streams (asynchronous, durable)
- **Agent → UI (Frontend):** WebSocket (real-time, bidirectional)
- **Human → System (API):** REST over HTTP (synchronous, simple)
- **NOT gRPC:** Complexity not justified for test automation latency requirements

WebSocket for Real-Time UI Communication

WebSocket Advantages for Agent Dashboards:[\[13\]](#) [\[14\]](#) [\[15\]](#) [\[16\]](#)

- **Full-duplex bidirectional** communication over single connection
- **Stateful:** Connection persists, maintains context

- **Sub-millisecond latency** for established connections
- **Perfect for:** Streaming agent outputs, progress updates, human-in-the-loop scenarios

Implementation Pattern: [\[14\]](#) [\[15\]](#)

```
@app.websocket("/ws/agent/dashboard")
async def agent_dashboard(websocket):
    await websocket.accept()

    # Subscribe to agent events (Redis Pub/Sub)
    pubsub = redis.pubsub()
    pubsub.subscribe('ui.agent.status')

    # Stream updates to UI
    async for message in pubsub.listen():
        await websocket.send_json({
            "type": "agent_status_update",
            "data": message['data']
        })
```

Use Cases for Test Automation: [\[15\]](#) [\[13\]](#)

- Live test generation progress (percentage complete, tests created)
- Agent status dashboard (active/idle/error states)
- Real-time coverage metrics visualization
- Human oversight: Pause/resume/interrupt agent workflows [\[15\]](#)

Not for Backend: WebSocket is for **UI communication only**. Use message bus for agent-to-agent coordination.

Modern Agent Communication Protocols: 2025-2026 Survey

Four Emerging Protocols [\[17\]](#) [\[18\]](#)

Recent research (published May 2025) surveyed next-generation agent communication standards: [\[18\]](#)

MCP (Model Context Protocol):

- Purpose: LLM-to-tool invocation (client-server)
- Architecture: JSON-RPC over HTTP
- Best for: Secure tool access, typed data exchange
- Limitation: Not peer-to-peer

ACP (Agent Communication Protocol):

- Purpose: General-purpose agent messaging
- Heritage: FIPA ACL/KQML modernized with HTTP

- Features: Performative-based (inform, request, propose)
- Best for: Logic-based and rule-based agents

A2A (Agent-to-Agent Protocol):

- Purpose: Peer-to-peer task delegation
- Architecture: Capability-based Agent Cards
- Best for: Enterprise collaborative workflows
- Key feature: Cross-ecosystem interoperability

ANP (Agent Network Protocol):

- Purpose: Decentralized agent marketplaces
- Features: Discovery, routing, health monitoring
- Identity: W3C DIDs (Decentralized Identifiers)
- Best for: Multi-organization agent networks

Recommendation: Don't Adopt Yet

Why Not?

1. **Emerging standards** (2025-2026)—not production-proven
2. **Limited tooling and community support**
3. **Overkill** for single-organization test automation system
4. **Adoption risk:** Standards may not consolidate

Learn concepts (performatives, capability-based routing) but **implement simplified version** with proven technologies (Redis Streams + JSON messages).

Future-proofing: Design message schemas to be protocol-agnostic. If A2A or ACP becomes dominant in 2027+, migration path exists via adapter layer.

Message Schema Design: Production-Grade Patterns

JSON Schema with Semantic Versioning

Design Principles: [\[19\]](#) [\[20\]](#) [\[21\]](#) [\[22\]](#)

```
{
  "$schema": "http://json-schema.org/draft-07/schema#",
  "title": "AgentMessage",
  "type": "object",
  "properties": {
    "schema_version": {
      "type": "string",
      "pattern": "\d+\.\d+\.\d+",
      "description": "Semantic version (MAJOR.MINOR.PATCH)"
    }
  }
}
```

```
"message_id": {
    "type": "string",
    "format": "uuid"
},
"conversation_id": {
    "type": "string",
    "format": "uuid",
    "description": "Thread/session identifier for tracing"
},
"sender_id": {
    "type": "string",
    "pattern": "^agent_[a-z_]+$"
},
"receiver_id": {
    "type": "string",
    "pattern": "^agent_[a-z_]+|\$\*\$",
    "description": "Target agent or * for broadcast"
},
"performative": {
    "type": "string",
    "enum": ["request", "inform", "propose", "accept", "reject", "query", "confirm"],
    "description": "FIPA-inspired speech act"
},
"content": {
    "type": "object",
    "properties": {
        "type": {
            "type": "string",
            "description": "Content discriminator (e.g., test_pattern_detected)"
        },
        "payload": {
            "type": "object"
        }
    },
    "required": ["type", "payload"]
},
"priority": {
    "type": "integer",
    "minimum": 0,
    "maximum": 10,
    "default": 5
},
"ttl": {
    "type": "integer",
    "minimum": 0,
    "description": "Time-to-live in seconds"
},
"timestamp": {
    "type": "string",
    "format": "date-time"
},
"reply_to": {
    "type": ["string", "null"],
    "format": "uuid"
},
"metadata": {
```

```

    "type": "object",
    "properties": {
        "retry_count": {"type": "integer", "default": 0},
        "trace_id": {"type": "string"}
    },
    "additionalProperties": true
}
},
"required": ["schema_version", "message_id", "conversation_id", "sender_id", "receiver_id"]
}

```

Key Design Decisions:

1. Simplified Performatives: 7 types (vs FIPA's 20+) [\[17\]](#) [\[23\]](#)

- Keeps cognitive overhead low
- Covers 95% of agent communication needs
- Performatives: request, inform, propose, accept, reject, query, confirm

2. Content Encapsulation: Two-level structure (type + payload)

- type discriminates content semantics (e.g., "test_pattern_detected")
- payload carries performative-specific data
- Enables schema validation per content type

3. Priority System: 0-10 scale for urgency

- 0-3: Low (background analytics)
- 4-6: Normal (standard test generation)
- 7-9: High (user-initiated requests)
- 10: Critical (system health, emergency shutdown)

4. TTL Mechanism: Auto-expire stale messages

- Prevents queue buildup from slow consumers
- Example: Status update TTL=30s (irrelevant after 30 seconds)

5. Extensible Metadata: Future-proofing

- additionalProperties: true allows new fields without schema migration
- Backward compatibility: Old agents ignore unknown metadata fields

Message Versioning Strategy [\[24\]](#) [\[25\]](#) [\[26\]](#)

Semantic Versioning (MAJOR.MINOR.PATCH): [\[24\]](#) [\[25\]](#)

- **MAJOR** (1.0.0 → 2.0.0): Breaking changes

- Remove required fields
- Change field types (string → integer)
- Rename fields

- **Migration required:** All agents must upgrade
- **MINOR** (1.0.0 → 1.1.0): Backward-compatible additions
 - Add optional fields
 - Add new performatives or content types
 - **No migration required:** Old agents ignore new fields
- **PATCH** (1.0.0 → 1.0.1): Bug fixes
 - Clarify descriptions
 - Fix validation rules
 - **No migration required:** Semantic-only changes

Version Support Policy:

- Agents **MUST** support **N-1 versions** (1 version backward)
- Agents **SHOULD** support **N+1 versions** (forward compatibility via ignoring unknown fields)
- **Deprecation window:** 3 months notice before removing old version support

Schema Registry:

- Central storage in PostgreSQL or Redis
- Version lookup: GET /schemas/AgentMessage/1.2.0
- Validation on send/receive: Reject messages with unsupported schema versions

Migration Example:

```

Version 1.0.0 → 1.1.0 (MINOR)
+ Added: "priority" field (default: 5)
+ Added: "ttl" field (default: 0)
Result: Old agents ignore these fields, continue working

Version 1.1.0 → 2.0.0 (MAJOR)
- Removed: "legacy_format" field
~ Changed: "timestamp" from Unix epoch to ISO 8601
Result: All agents must upgrade within deprecation window
  
```

Error Handling & Retry Architecture

Exponential Backoff with Jitter

Algorithm: [\[27\]](#) [\[28\]](#) [\[29\]](#)

Exponential backoff retry pattern with dead letter queue for failed agent messages

Implementation: [\[27\]](#) [\[28\]](#)

```

def calculate_retry_delay(attempt: int, base_delay: float = 1.0, max_delay: float = 60.0)
    """Calculate exponential backoff with jitter"""
  
```

```

# Exponential: 1s, 2s, 4s, 8s, 16s, 32s, 60s (capped)
delay = min(base_delay * (2 ** attempt), max_delay)

# Add 10% jitter to prevent thundering herd
jitter = random.uniform(0, delay * 0.1)

return delay + jitter

```

Retry Schedule:

- Attempt 1: 1.0s ± 0.1s
- Attempt 2: 2.0s ± 0.2s
- Attempt 3: 4.0s ± 0.4s
- Attempt 4: 8.0s ± 0.8s
- Attempt 5: 16.0s ± 1.6s
- **After 5 failures:** Move to Dead Letter Queue

Why Jitter Matters: [\[28\]](#) [\[27\]](#)

- **Problem:** Multiple agents fail simultaneously (e.g., downstream service outage)
- **Without jitter:** All agents retry at identical intervals → thundering herd
- **With jitter:** Retry attempts spread over time window (10% variance)
- **Result:** Reduced load spikes on recovering services

Queue-Based Retry Pattern

Three-Step Process: [\[27\]](#)

1. **Delete:** Remove original message from processing queue
2. **Calculate:** Compute exponential backoff delay
3. **Requeue:** Send to retry queue with delay timestamp + incremented retry count

Advantages Over In-Process Retry: [\[27\]](#)

- **No memory leaks:** Queue manages message lifecycle
- **Distributed-friendly:** Works across multiple agent instances (no shared state)
- **Durable:** Survives process crashes (persisted in Redis)
- **Cost-effective:** Pay only for messages processed (vs idle timers)
- **Observable:** Queue depth metrics show retry backlog

Implementation with Redis Streams:

```

async def handle_message_failure(
    message: dict,
    error: Exception,
    redis: Redis
):

```

```

retry_count = message.get("metadata", {}).get("retry_count", 0)

# Check max retries
if retry_count >= MAX_RETRIES:
    await move_to_dlq(message, error, redis)
    return

# Calculate backoff
delay = calculate_retry_delay(retry_count)
retry_at = datetime.utcnow() + timedelta(seconds=delay)

# Update metadata
message["metadata"]["retry_count"] = retry_count + 1
message["metadata"]["retry_at"] = retry_at.isoformat()
message["metadata"]["last_error"] = str(error)

# Requeue with delay
await redis.xadd(
    "agent:retry_queue",
    {
        "message": json.dumps(message),
        "process_after": retry_at.isoformat()
    }
)

logger.info(f"Message {message['message_id']} requeued, attempt {retry_count + 1}, de

```

Error Classification:[\[30\]](#) [\[27\]](#)

- **Transient errors (RETRY)**: Network timeout, rate limit, service unavailable (503)
- **Permanent errors (FAIL FAST)**: Schema validation, authorization (401/403), logical errors
- **Unknown errors (RETRY WITH CAUTION)**: Log for investigation, max 2 attempts

Dead Letter Queue (DLQ)

Purpose:[\[31\]](#) [\[32\]](#) [\[33\]](#)

- Store messages that **exhausted max retries**
- **Audit trail**: Prevent message loss, enable forensics
- **Unblock processing**: Failed messages don't clog healthy flow
- **Manual recovery**: Human intervention point

DLQ Message Schema:

```
{
  "dlq_id": "uuid",
  "original_message": {...},
  "failure_reason": "Connection timeout to analysis service",
  "error_type": "NetworkError",
  "retry_history": [
    {"attempt": 1, "timestamp": "2026-01-17T10:00:00Z", "error": "Connection timeout"},
    {"attempt": 2, "timestamp": "2026-01-17T10:00:02Z", "error": "Connection timeout"}
  ]
}
```

```

        {"attempt": 3, "timestamp": "2026-01-17T10:00:06Z", "error": "Connection timeout"},  

        {"attempt": 4, "timestamp": "2026-01-17T10:00:14Z", "error": "Connection timeout"},  

        {"attempt": 5, "timestamp": "2026-01-17T10:00:30Z", "error": "Connection timeout"}  

    ],  

    "final_retry_count": 5,  

    "moved_to_dlq_at": "2026-01-17T10:00:30Z",  

    "metadata": {  

        "sender_id": "agent_orchestration",  

        "receiver_id": "agent_analysis",  

        "conversation_id": "abc-123"  

    }  

}

```

DLQ Monitoring & Alerting:[\[33\]](#) [\[31\]](#)

- **Alert 1:** DLQ count > 100 messages
- **Alert 2:** DLQ growth rate > 50/hour
- **Alert 3:** Message age in DLQ > 1 hour
- **Alert 4:** Specific error type spike (e.g., 20+ "Connection timeout" in 5 min)

Recovery Strategies:[\[31\]](#) [\[33\]](#)

1. **Automatic Replay:** Once downstream service recovers, replay DLQ in order
2. **Manual Fix + Replay:** Human corrects data/config, triggers replay
3. **Transform + Replay:** Fix message schema, replay with corrections
4. **Discard:** For unrecoverable messages (log for audit, send alert)

Message Routing & Event-Driven Patterns

Three Communication Patterns[\[34\]](#) [\[35\]](#)

1. Point-to-Point (Queue):[\[34\]](#) [\[35\]](#)

- **Semantics:** One message → One consumer
- **Implementation:** Redis Streams consumer groups
- **Use for:** Task allocation, command execution
- **Example:** Orchestration Agent → Specific Analysis Agent

2. Publish-Subscribe (Topic):[\[35\]](#) [\[34\]](#)

- **Semantics:** One message → Many consumers
- **Implementation:** Redis Pub/Sub channels
- **Use for:** Event broadcasting, status notifications
- **Example:** Test execution completed → All interested agents

3. Request-Reply (Synchronous):[\[34\]](#) [\[35\]](#)

- **Semantics:** Sender waits for response

- **Implementation:** Correlation ID + reply-to address
- **Use for:** Queries, commands requiring confirmation
- **Example:** UI requests test count → Reporting Agent responds

Event Message Types ^[34]

Three Levels of Data Transfer: ^[34]

1. Notification Only: Key + link

```
{
  "performative": "inform",
  "content": {
    "type": "test_created",
    "payload": {
      "test_id": "test-789",
      "link": "/api/tests/test-789"
    }
  }
}
```

- **Pros:** Minimal data transfer, receiver fetches if interested
- **Cons:** Extra roundtrip for details

2. Notification + Metadata: Key + changed fields

```
{
  "performative": "inform",
  "content": {
    "type": "coverage_changed",
    "payload": {
      "project_id": "proj-456",
      "old_coverage": 0.75,
      "new_coverage": 0.82,
      "change": "+0.07"
    }
  }
}
```

- **Pros:** Enough data to decide action, no fetch needed
- **Cons:** More data in message

3. Event Carried State Transfer: Full data payload

```
{
  "performative": "inform",
  "content": {
    "type": "pattern_detected",
    "payload": {
      "pattern_id": "duplicate_assertion",
      "confidence": 0.87,
      "test_cases": ["test_user_login", "test_admin_login"],
      "suggestion": "Refactor into parameterized test",
    }
  }
}
```

```

        "code_snippet": "..."
    }
}
}
```

- **Pros:** Receiver has all data immediately, no fetch
- **Cons:** Larger messages, potential data duplication

Recommendation: Use **Notification + Metadata** as default, full payload only when needed for agent decision-making.

Message Flow Example

Test Generation Workflow:

1. UI → Orchestration Agent (HTTP POST, synchronous)


```
POST /api/agent/generate-tests
Body: {"repo_url": "...", "branch": "main", "target_coverage": 0.85}
```
2. Orchestration Agent → Observation Agent (Redis Streams, Point-to-Point)


```
Stream: agent:observation:inbox
Message: {
    "performative": "request",
    "content": {
        "type": "analyze_test_code",
        "payload": {"repo_url": "...", "branch": "main"}
    }
}
```
3. Observation Agent → Orchestration Agent (Redis Streams, Reply)


```
Stream: agent:orchestration:inbox
Message: {
    "performative": "inform",
    "reply_to": "<message_id_from_step_2>",
    "content": {
        "type": "analysis_complete",
        "payload": {"patterns": [...], "coverage": 0.75}
    }
}
```
4. Orchestration Agent → ALL Agents (Redis Pub/Sub, Broadcast)


```
Channel: events.system.broadcast
Message: {
    "performative": "inform",
    "receiver_id": "*",
    "content": {
        "type": "test_generation_started",
        "payload": {"session_id": "session-abc"}
    }
}
```
5. Orchestration Agent → Evolution Agent (Redis Streams, Point-to-Point)


```
Stream: agent:evolution:inbox
Message: {
```

```

        "performative": "request",
        "content": {
            "type": "generate_tests",
            "payload": {"patterns": [...], "target_coverage": 0.85}
        }
    }

6. Evolution Agent → UI (WebSocket, Push)
WebSocket: /ws/agent/dashboard
Message: {
    "type": "generation_progress",
    "tests_generated": 5,
    "estimated_total": 12,
    "progress": 0.42
}

7. Evolution Agent → Orchestration Agent (Redis Streams, Reply)
Stream: agent:orchestration:inbox
Message: {
    "performative": "inform",
    "reply_to": "<message_id_from_step_5>",
    "content": {
        "type": "tests_generated",
        "payload": {"test_count": 12, "tests": [...]}}
    }
}

8. Orchestration Agent → UI (HTTP Response, synchronous)
Response: {
    "status": "complete",
    "tests_generated": 12,
    "coverage_improvement": 0.10,
    "session_id": "session-abc"
}

```

Production Implementation Roadmap

Redis Infrastructure Setup

Cluster Configuration:

- **Redis Cluster:** 3+ nodes with replication
- **Redis Sentinel:** Automatic failover monitoring
- **Persistence:** AOF (Append-Only File) + RDB snapshots
- **Backup:** Daily snapshots to S3/object storage

Message Streams (Redis Streams):

- agent:orchestration:inbox - Commands to orchestration agent
- agent:observation:inbox - Commands to observation agent
- agent:requirements:inbox - Commands to requirements agent

- `agent:analysis:inbox` - Commands to analysis agent
- `agent:evolution:inbox` - Commands to evolution agent
- `agent:reporting:inbox` - Commands to reporting agent
- `agent:retry_queue` - Messages awaiting retry with delay
- `agent:dlq` - Dead letter queue for failed messages

Pub/Sub Channels (Redis Pub/Sub):

- `events.test.created` - Test creation notifications
- `events.test.executed` - Test execution notifications
- `events.pattern.detected` - Pattern detection notifications
- `events.coverage.changed` - Coverage change notifications
- `events.system.broadcast` - System-wide announcements
- `ui.agent.status` - Agent status updates for dashboard

Performance Targets

Metric	Target	Measurement
Redis Streams latency	<5ms P99	Time from XADD to XREAD
Message processing	<100ms P95	Agent receives → ACK sent
End-to-end task	<30s P95	Request → Response
Queue depth	<5,000 messages	Alert if exceeded
DLQ size	<100 messages	Alert if exceeded
Message age	<15 min P99	Alert on stale messages

Observability & Monitoring

Key Metrics:

- **Throughput:** Messages/sec per stream
- **Latency:** Processing time distribution (P50, P95, P99)
- **Queue Depth:** Backlog size per agent inbox
- **Error Rate:** Failed messages / total messages
- **Retry Rate:** Messages retried / total messages
- **DLQ Size:** Messages in dead letter queue

Tracing:

- **Correlation ID:** `conversation_id` propagates across all messages
- **OpenTelemetry:** Span per agent processing step
- **Distributed Tracing:** Full request flow from UI → Orchestrator → Agents → UI

Alerting:

- Queue depth > 5,000 (capacity warning)
- DLQ size > 100 (failure spike)
- Message age > 15 min (processing stall)
- Error rate > 10% (system degradation)
- Specific agent down > 5 min (availability issue)

Final Architecture Recommendation

Communication Stack

Layer 1: Agent-to-Agent Coordination (Backend)

- **Primary:** Redis Streams
- **Pattern:** Point-to-Point via consumer groups
- **Message Type:** Commands, queries, responses
- **Delivery:** At-least-once with acknowledgment
- **Latency:** 1-5ms
- **Rationale:** Guaranteed delivery, persistence for debugging, consumer groups for scaling

Layer 2: Event Broadcasting (Backend)

- **Primary:** Redis Pub/Sub
- **Pattern:** Publish-Subscribe via channels
- **Message Type:** Notifications, status updates
- **Delivery:** At-most-once (fire-and-forget)
- **Latency:** <1ms
- **Rationale:** Real-time, transient events where loss acceptable

Layer 3: UI Communication (Frontend)

- **Primary:** WebSocket
- **Pattern:** Bidirectional streaming
- **Message Type:** Real-time updates, progress notifications
- **Delivery:** At-most-once
- **Latency:** <2ms
- **Rationale:** Native browser support, stateful connection for dashboard

Layer 4: External API (Human-facing)

- **Primary:** REST over HTTP
- **Pattern:** Request-Reply

- **Message Type:** User commands, queries
- **Delivery:** Synchronous
- **Latency:** <500ms (backend processing time)
- **Rationale:** Simplicity, ubiquitous tooling, easy debugging

Why This Hybrid Approach Wins

1. **Best Tool for Each Job:** Streams for durability, Pub/Sub for speed, WebSocket for UI, REST for humans
2. **Operational Simplicity:** All Redis-based (single technology), well-understood by DevOps teams
3. **Battle-Tested:** Millions of deployments, mature ecosystem
4. **Cost-Effective:** Open-source Redis vs commercial message brokers
5. **Performance:** Sub-5ms latency for agent coordination, sub-1ms for notifications
6. **Scalability:** 1M+ messages/sec throughput, horizontal scaling with consumer groups

What We're NOT Using (and Why)

Kafka: Too heavyweight for agent messaging (10-50ms latency, complex operations)

RabbitMQ: Slower than Redis (5-20ms latency, 50-60K msg/sec vs 1M+)

gRPC: Implementation complexity not justified (45min setup vs minutes for Redis)

FIPA ACL/KQML: Too heavyweight for modern systems (ontology overhead)

A2A/ACP/APN: Emerging protocols, not production-proven in 2026

Key Takeaways

Critical Decisions

1. **Redis Streams (NOT Pub/Sub)** for agent task coordination
 - Persistence, consumer groups, exactly-once semantics
 - Only 1-5ms latency penalty vs Pub/Sub's <1ms
 - Production requirement: Message history for debugging
2. **Simplified Performatives** (7 types, not FIPA's 20+)
 - request, inform, propose, accept, reject, query, confirm
 - 95% of communication needs with 5% of complexity
3. **Exponential Backoff with Jitter** for retries
 - 10% jitter prevents thundering herd
 - Max 5 retries: 1s, 2s, 4s, 8s, 16s
 - Dead Letter Queue after max retries
4. **Semantic Versioning** for message schemas

- MAJOR.MINOR.PATCH
- Agents support N-1 and N+1 versions (backward + forward compatibility)

5. Hybrid Message Patterns

- Point-to-Point: Task allocation (Redis Streams)
- Pub/Sub: Event broadcasting (Redis Pub/Sub)
- Request-Reply: Synchronous queries (HTTP REST)
- Streaming: Real-time UI (WebSocket)

Implementation Priorities (Sprint 7)

Week 1:

- [] Setup Redis Cluster (3 nodes, replication, Sentinel)
- [] Define message schemas (JSON Schema Draft 7)
- [] Implement message validation middleware
- [] Create base sender/receiver classes

Week 2:

- [] Implement retry logic with exponential backoff
- [] Setup Dead Letter Queue stream and monitoring
- [] Create WebSocket endpoint for UI dashboard
- [] Write integration tests for message passing
- [] Setup observability (OpenTelemetry, metrics)

Success Criteria

Functional:

- ✓ Message passing between 2 agents <5ms latency
- ✓ Retry mechanism tested (5 attempts, correct delays)
- ✓ DLQ receives messages after max retries
- ✓ WebSocket streams real-time updates to UI
- ✓ Schema validation rejects malformed messages

Non-Functional:

- ✓ Redis Cluster survives 1 node failure (HA test)
- ✓ Consumer groups distribute work across 3 agent instances
- ✓ Message replay works after simulated crash
- ✓ Observability: All messages traced with conversation_id
- ✓ Performance: 10,000 msg/sec throughput sustained

This architecture balances **performance** (sub-5ms latency), **reliability** (guaranteed delivery with retries), **scalability** (consumer groups for horizontal scaling), and **operational simplicity** (single Redis technology stack). It's production-ready, battle-tested, and right-sized for Phase 3 multi-agent test automation.

**

1. <https://www.geeksforgeeks.org/system-design/difference-between-redis-pub-sub-vs-redis-streams/>
2. <https://stackoverflow.com/questions/59540563/what-are-the-main-differences-between-redis-valkey-pub-sub-and-streams>
3. <https://dev.to/lovestaco/redis-pubsub-vs-redis-streams-a-dev-friendly-comparison-39hm>
4. <https://www.linkedin.com/pulse/redis-pubsub-vs-streams-osama-ahmed>
5. <https://leapcell.io/blog/redis-messaging-showdown-pub-sub-vs-streams-for-event-driven-architectures>
6. <https://www.linkedin.com/pulse/redis-streams-vs-pubsub-performance-perspective-ykr9c>
7. <https://www.index.dev/skill-vs-skill/backend-grpc-vs-rest-vs-graphql>
8. <https://www.digiratina.com/blogs/rest-vs-grpc-a-real-world-performance-experiment/>
9. <https://l3montree.com/blog/performance-comparison-rest-vs-grpc-vs-asynchronous-communication>
10. <https://blog.dreamfactory.com/grpc-vs-rest-how-does-grpc-compile-with-traditional-rest-apis>
11. <https://oneuptime.com/blog/post/2026-01-08-grpc-vs-rest-api-comparison/view>
12. <https://eluminoustechnologies.com/blog/grpc-vs-rest/>
13. <https://blog.algomaster.io/p/websocket-use-cases-system-design>
14. <https://developers.cloudflare.com/agents/api-reference/websockets/>
15. <https://aws.amazon.com/cn/blogs/china/fast-fashion-e-commerce-agent-design-8-websocket-voice-system/>
16. <https://dev.to/giwajossy/understanding-websocket-real-time-communication-made-easy-5904>
17. <https://k21academy.com/ai-ml/agentic-ai/agentic-ai-protocols-comparison/>
18. <https://arxiv.org/abs/2505.02279>
19. <https://json-schema.org/learn/getting-started-step-by-step>
20. <https://docs.solace.com/Schema-Registry/schema-registry-best-practices.htm>
21. <https://www.tencentcloud.com/techpedia/128070>
22. <https://www.ibm.com/docs/en/app-connect/12.0.x?topic=schema-json-requirements-message-maps>
23. <https://smythos.com/developers/agent-development/fipa-agent-communication-language/>
24. <https://www.designgurus.io/answers/detail/how-do-you-handle-versioning-in-microservices>
25. <https://www.serverion.com/uncategorized/versioning-strategies-for-microservices-schemas/>
26. <https://stackoverflow.com/questions/39485459/microservice-versioning>
27. <https://dev.to/andreparsis/queue-based-exponential-backoff-a-resilient-retry-pattern-for-distributed-systems-37f3>
28. <https://www.hackerone.com/blog/retrying-and-exponential-backoff-smart-strategies-robust-software>
29. https://en.wikipedia.org/wiki/Exponential_backoff
30. <https://techholding.co/blog/understanding-message-queues-and-retry-architectures>

31. <https://dev.to/mehmetakar/dead-letter-queue-3mj6>
32. <https://www.geeksforgeeks.org/system-design/dead-letter-queue-system-design/>
33. <https://ibm-cloud-architecture.github.io/refarch-eda/patterns/dlq/>
34. <https://solace.com/event-driven-architecture-patterns/>
35. <https://solace.com/blog/messaging-patterns-for-event-driven-microservices/>
36. [Phase3-Developer-A-Research-Plan.md](#)
37. https://dev.to/sizan_mahmud0_e7c3fd0cb68/the-complete-guide-to-api-types-in-2026-rest-graphql-grpc-soap-and-beyond-191
38. <https://www.wallarm.com/what/grpc-vs-rest-comparing-key-api-designs-and-deciding-which-one-is-best>
39. <https://smythos.com/developers/agent-development/agent-communication-languages-and-protocols-comparison/>
40. <https://w3c.github.io/json-ld-bp/>
41. <https://json-schema.org/UnderstandingJSONSchema.pdf>
42. https://unece.org/sites/default/files/2023-11/API-TECH-SPEC_JSON_Schema_NDR_version1p0.pdf
43. <https://en.wikipedia.org/wiki/Event-driven.messaging>
44. <https://json-schema.org/understanding-json-schema/basics>
45. <https://www.geeksforgeeks.org/system-design/message-driven-architecture-vs-event-driven-architecture/>
46. <https://dev.to/lovestaco/choosing-the-right-messaging-tool-redis-streams-redis-pubsub-kafka-and-more-577a>
47. https://www.reddit.com/r/softwarearchitecture/comments/1nw3e1h/stop_confusing_redis_pubsub_with_streams/