



Please act as professional software architect designer to research on Day 3 research topic, challenge the finds with industrial best practice and provide the best solution after internal debate.

Based on my comprehensive Day 3 research, I've completed an in-depth analysis of agent orchestration patterns. Here's my professional architectural assessment after challenging findings with industrial best practices:

Agent Orchestration Patterns: Professional Architecture Analysis

Executive Summary

After analyzing 45+ authoritative sources covering orchestration vs choreography, Netflix Conductor, Temporal workflows, service discovery, and distributed deadlock detection, I recommend a **hybrid orchestration architecture** that combines centralized workflow control with decentralized event broadcasting. This pattern, proven in production at Netflix, Uber, and Airbnb, balances predictability with scalability.

Orchestration vs Choreography: The False Dichotomy

Industry Reality: Hybrid is Standard Practice

The debate is often framed as **orchestration vs choreography**, but this is a false choice. **Production systems use both patterns.** [\[1\]](#) [\[2\]](#) [\[3\]](#)

Orchestration vs Choreography: Key architectural differences for multi-agent system coordination

Why Pure Patterns Fail

Pure Orchestration Failures: [\[4\]](#) [\[5\]](#) [\[3\]](#)

- **Bottleneck:** All coordination flows through orchestrator (single point of saturation)
- **Tight Coupling:** Services become orchestrator-dependent
- **Scaling Ceiling:** Orchestrator CPU becomes limiting factor at 1000+ req/sec

- **Single Point of Failure:** Despite HA, coordination still centralized

Real-World Data: Orchestration reliability suffers without built-in transaction management, while choreography's point-to-point communication creates complex fault tolerance scenarios.^[3]

Pure Choreography Failures:^[5] ^[4]

- **State Fragmentation:** No single view of workflow state
- **Debugging Nightmare:** Event chains span 5-10 services (distributed tracing required)
- **Coordination Overhead:** Grows **non-linearly** with service count^[6]
- **Business Rule Enforcement:** No central authority to validate constraints

Real-World Example: Choreography complexity becomes unmanageable when "the number of events produced and consumed increases... managing these event chains may be difficult, especially for developers attempting to determine what caused a certain action."^[5]

The Hybrid Solution

Architecture Pattern:^[1] ^[2]

Hybrid orchestration architecture combining centralized control (orchestrator) with decentralized event broadcasting (pub/sub)

Critical Path (Orchestrated):

- Test generation workflow: Observe → Requirements → Analysis → Evolution → Report
- Orchestrator maintains **state machine**
- **Synchronous coordination** for workflow steps
- **Retry logic** and **error handling** centralized
- **Transaction-like semantics** (all-or-nothing for critical operations)

Side Effects (Choreographed):

- Status updates (agent health, progress bars)
- Metrics collection (coverage, execution time)
- Notifications (test completed, errors)
- **Asynchronous events** via Pub/Sub
- **Fire-and-forget** (loss acceptable)

Why This Works:^[1] ^[2]

1. **Flow distributed and protected** against single point of failure
2. **Services decoupled to an extent** (not completely)
3. **Optimized results** by tailoring each approach to specific needs
4. **Blends best of both worlds**

Netflix Conductor: Production Orchestration at Scale

Why Netflix Built Conductor^{[7] [8] [9]}

The Problem:^[8]

Traditional microservices at Netflix orchestrated workflows ad-hoc using:

- Pub/Sub for events
- Direct REST calls between services
- Database for state management

As microservices grew from dozens to hundreds, "**getting visibility into these distributed workflows becomes difficult without a central orchestrator**".^[8]

Conductor Architecture

Core Components:^{[7] [8]}

1. Workflow Definition (JSON DSL):^[7]

```
{  
    "name": "test_generation_workflow",  
    "version": 1,  
    "tasks": [  
        {  
            "name": "observe_code",  
            "taskReferenceName": "observe",  
            "type": "SIMPLE",  
            "inputParameters": {...}  
        },  
        {  
            "name": "analyze_patterns",  
            "taskReferenceName": "analyze",  
            "type": "SIMPLE",  
            "inputParameters": {  
                "patterns": "${observe.output.patterns}"  
            }  
        },  
        {  
            "name": "generate_tests",  
            "taskReferenceName": "evolve",  
            "type": "SIMPLE"  
        }  
    ],  
    "timeoutSeconds": 600,  
    "retryLogic": "FIXED",  
    "retryDelaySeconds": 60,  
    "retryCount": 3  
}
```

2. Task Workers (Polling Model):^{[8] [9]}

- Workers **poll** for pending tasks (backpressure handling)
- Idempotent stateless functions
- No need for worker discovery (workers find Conductor)
- Auto-scalability based on queue depth

Advantage over Push Model: "The polling model allows us to handle backpressure on the workers and provide auto-scalability based on the queue depth when possible." [\[8\]](#)

3. Decider (State Evaluator): [\[8\]](#)

- Evaluates workflow state against blueprint
- Schedules next tasks
- Handles conditional logic (if/else, loops, parallel execution)

4. Queue Abstraction: [\[8\]](#)

- Backend: Kafka, SQS, Redis
- Workers abstracted from queue implementation
- Conductor manages queue lifecycle

Conductor Features Applicable to Test Automation

Pause/Resume/Restart: [\[7\]](#) [\[8\]](#)

- Critical for **human-in-the-loop** scenarios
- User reviews generated tests, approves/rejects
- Workflow pauses, awaits human input, resumes

Workflow Visualization: [\[7\]](#)

- Real-time DAG showing task status
- Identify bottlenecks visually
- Debug failures (see exact task that failed)

Retry Policies: [\[7\]](#)

```
{
  "retryLogic": "EXPONENTIAL_BACKOFF",
  "retryDelaySeconds": 60,
  "maxRetryDelaySeconds": 300,
  "retryCount": 5,
  "backoffScaleFactor": 2.0
}
```

System Tasks: [\[8\]](#) [\[7\]](#)

- **Fork/Join:** Parallel agent execution
- **Decision:** Conditional branching (if coverage <0.8, generate more tests)

- **Sub-workflow:** Hierarchical workflows (test generation → test execution)
- **HTTP Task:** REST calls without custom workers

Scalability: [8] [7]

- **Millions of concurrent workflows** at Netflix production
- Proven at scale

When NOT to Use Conductor

Overhead for Simple Use Cases: [9]

- Requires Conductor server infrastructure
- Queue backend (Kafka/SQS)
- Metadata storage (database)
- Operational complexity

Alternative: **LangGraph** provides similar orchestration features with **minimal infrastructure** (just Redis + PostgreSQL for checkpointing).

Verdict: Conductor patterns are excellent reference, but **LangGraph** is right-sized for test automation.

Temporal: Durable Workflow Engine

Core Innovation: Workflows as Deterministic State Machines [10] [11]

Key Principle: Workflows are **deterministic code** that can be replayed from history to reconstruct state. [10]

Event Sourcing Architecture: [10]

1. Every state transition persisted as event
2. Workflow execution = replay all events
3. Crash recovery = replay from last checkpoint
4. Time travel testing = replay with fast-forwarded clock [11]

Example: Time Travel Testing: [11]

```
@pytest.mark.asyncio
async def test_pickup_timeout():
    # Create order
    order = await workflow.create_order(...)

    # Fast-forward 1 day to force delivery timeout
    testEnv.sleep(Duration.ofDays(1))

    # Workflow completes timeout logic instantly
    result = await workflow.get_result()
```

```
assert result.status == "cancelled"
```

This is **powerful** for testing multi-day workflows without waiting.[\[11\]](#)

Temporal Architecture[\[10\]](#)

History Component: Manages state transitions of individual workflows

Transfer Queues: Transactional task creation (atomically commit to DB + queue)

Matching Component: Delivers tasks to workers, queues poll requests

Worker Component: System workflows (batch operations, database scans)

Advantages over Conductor:[\[10\]](#)

- **Code as workflow** (not JSON DSL) - full programming language flexibility
- **Built-in deterministic testing** (time travel)
- **Long-running workflows** (months/years, not hours/days)

Disadvantages:

- **Steeper learning curve** (determinism constraints)
- **Deterministic code required** (no random(), no time.now(), must use workflow.random()/workflow.now())
- **LLM calls are non-deterministic** (problem for AI agents)

Why Temporal Doesn't Fit Test Automation

Critical Issue: LLM outputs are **non-deterministic**.[\[10\]](#)

Temporal requires deterministic execution for replay. But:

- GPT-4 generates different tests each run (temperature >0)
- Pattern detection confidence scores vary
- Embedding similarity changes

Workaround: Wrap LLM calls in Activities (non-deterministic allowed), but this **negates Temporal's replay advantage**.

Verdict: Temporal's deterministic state machine is **incompatible** with LLM non-determinism. Patterns (checkpointing, retry) are valuable, but **LangGraph better fit** for AI agents.

Task Allocation: Contract Net Protocol

Classic CNP (1980) [\[12\]](#) [\[13\]](#)

Contract Net Protocol for dynamic task allocation in multi-agent systems

Modern Evolution: Agent Contracts (2026) [\[14\]](#)

Enhanced with: [\[14\]](#)

- **Resource constraints:** Token budgets (`max_tokens: 10000`), time limits (`deadline: 2026-01-18T12:00Z`)
- **Temporal boundaries:** SLA enforcement
- **Success criteria:** Coverage targets (`min_coverage: 0.80`), quality thresholds
- **Hierarchical delegation:** Winner can sub-contract to specialists
- **Budget conservation:** Track token usage, prevent cost overruns

Production Implementation:

```
# Orchestrator announces
cfp = {
    "task_id": "gen_tests_user_service",
    "task_type": "test_generation",
    "requirements": {"target_class": "UserService", "coverage": 0.85},
    "constraints": {"max_tokens": 10000, "deadline_seconds": 300},
    "success_criteria": {"min_coverage": 0.80}
}

# Evolution Agent bids
bid = {
    "agent_id": "agent_evolution_1",
    "confidence": 0.87, # 87% confident can meet criteria
    "estimated_tokens": 8500,
    "estimated_time_seconds": 180,
    "current_load": 3 # 3 active tasks
}

# Orchestrator awards
winner = max(bids, key=lambda b: b["confidence"] / (1 + b["current_load"]))
# Prefer high confidence + low load
```

When to Use CNP

✓ Use CNP When:

- Multiple agents capable of same task (2+ Evolution Agents)
- Dynamic load balancing needed (some agents overloaded)
- Resource awareness critical (token budgets, cost limits)
- Quality optimization (select best agent, not just any available)

✗ Skip CNP When:

- Only one agent per type (direct assignment faster)
- Latency-critical (bidding adds 100-500ms)
- Simple round-robin sufficient (complexity not justified)

Phased Approach:

- **Phase 1** (Sprints 7-8): Direct assignment (1 agent per type)
- **Phase 2** (Sprints 9-10): Weighted round robin (multiple instances)
- **Phase 3** (Sprints 11-12): Contract Net Protocol (optimal allocation)

Agent Discovery: Service Registry Pattern

Modern Agent Registry Architecture [\[15\]](#) [\[16\]](#) [\[17\]](#) [\[18\]](#)

Components: [\[15\]](#) [\[17\]](#) [\[18\]](#)

1. Central Registry (Redis):

```
# Agent Card stored in registry
{
  "agent_id": "agent_evolution_1",
  "agent_type": "evolution",
  "capabilities": [
    {"name": "test_generation", "version": "1.0"},
    {"name": "mutation_testing", "version": "1.1"}
  ],
  "endpoints": {"inbox_stream": "agent:evolution_1:inbox"},
  "health": {
    "status": "healthy",
    "last_heartbeat": "2026-01-18T10:30:00Z",
    "cpu_usage": 0.45,
    "active_tasks": 3
  },
  "constraints": {"max_concurrent_tasks": 5, "rate_limit": "100/hour"}
}
```

2. Self-Registration Pattern: [\[19\]](#)

- Agent registers on startup
- Sends heartbeat every 30 seconds
- Auto-expires if heartbeat >90s late
- De-registers on graceful shutdown

3. Discovery Patterns: [\[16\]](#) [\[17\]](#)

Client-Side Discovery: [\[16\]](#)

```
# Client queries registry
agents = await registry.find_agents({"capability": "test_generation", "status": "healthy"}
```

```

# Client selects (load balancing logic in client)
selected = min(agents, key=lambda a: a["health"]["active_tasks"])
# Client sends message directly
await send_to_agent(selected["endpoints"]["inbox_stream"], task)

```

Server-Side Discovery:[\[16\]](#)

```

# Client sends to router
await send_to_router("test_generation", task)
# Router queries registry, selects agent, forwards
# Client doesn't need load balancing logic

```

Recommendation: Client-side discovery for test automation (orchestrator has LB logic, agents are simple).

MCP-Based Discovery (2026 Emerging Pattern)[\[17\]](#)

Model Context Protocol enables **runtime capability enumeration**:[\[17\]](#)

```

# Agent exposes MCP server
mcp_server = AgentMCPServer(agent)

# Client queries capabilities dynamically
capabilities = await mcp_server.list_tools()
# Returns: ["analyze_code", "detect_patterns", "suggest_refactorings"]

# Client invokes tool
result = await mcp_server.invoke_tool("analyze_code", {"repo_url": "..."})

```

Advantages:[\[17\]](#)

- **Dynamic discovery:** Capabilities change at runtime (new models added)
- **No static registration:** Agents announce capabilities when queried
- **Real-time adaptation:** If agent adds new capability, immediately discoverable

Challenges:

- **Cutting-edge** (2025-2026 standard)
- **Limited tooling** vs REST
- **Overkill** for static capability set (test automation agents don't change often)

Verdict: MCP is **future-proofing pattern**, but **not needed** for Phase 3. Use **simple REST-based registry** with Agent Cards.

Graceful Shutdown & Health Monitoring

Graceful Shutdown (Production Pattern) [\[20\]](#) [\[21\]](#) [\[22\]](#)

Five Steps: [\[20\]](#) [\[21\]](#)

1. Stop Accepting New Requests:

```
self.accepting_requests = False
await registry.update_status(self.id, "shutting_down")
```

2. Complete In-Flight Requests (with timeout):

```
shutdown_timeout = 300 # 5 minutes
while self.active_tasks > 0 and not timeout_expired:
    await asyncio.sleep(1)
```

3. Persist State (checkpoint):

```
await self.save_checkpoint({
    "pending_tasks": list(self.task_queue),
    "partial_results": {...}
})
```

4. Release Resources:

```
await self.redis.close()
await self.message_queue.disconnect()
await self.release_all_locks()
```

5. De-Register:

```
await registry.deregister(self.id)
```

Kubernetes Integration: [\[20\]](#)

```
lifecycle:
  preStop:
    exec:
      command: ["curl", "-X", "POST", "http://localhost:8000/shutdown"]
terminationGracePeriodSeconds: 300
```

Critical: Kubernetes sends **SIGTERM**, waits `terminationGracePeriodSeconds` (300s), then sends **SIGKILL**. Agent **must** complete shutdown within 5 minutes or be force-killed.

Health Monitoring Patterns [\[23\]](#) [\[24\]](#) [\[25\]](#)

Three Health Check Types: [\[24\]](#)

Liveness (Is agent alive?):

```
GET /health/agent_evolution_1
→ {"status": "healthy", "uptime_seconds": 3600}
```

Readiness (Can agent accept work?):

```
GET /ready/agent_evolution_1
→ {"status": "ready"} if active_tasks < max_tasks else 503
```

Deep (All dependencies healthy?):

```
GET /health/agent_evolution_1/deep
→ {
    "status": "healthy",
    "checks": {
        "redis": true,
        "llm_api": true,
        "vector_db": true
    }
}
```

Monitoring Metrics: [\[23\]](#) [\[25\]](#) [\[24\]](#)

- **Availability:** Uptime % (target: 99.5%)
- **Latency:** P95 task execution time (<5s)
- **Error Rate:** Failed tasks % (<1%)
- **Queue Depth:** Pending tasks (<1000)
- **Resource Usage:** CPU (<80%), Memory (<80%)

Alerting: [\[24\]](#)

- **Critical:** Agent down >5 min, error rate >10%, queue >5000
- **Warning:** P95 latency >10s, CPU >80%, memory >80%

Deadlock Detection & Prevention

The Distributed Deadlock Problem [\[26\]](#) [\[27\]](#) [\[28\]](#)

Scenario: Agents wait for resources held by each other, forming circular dependency. [\[27\]](#)

Example:

```
Agent A holds Lock1, waits for Lock2  
Agent B holds Lock2, waits for Lock1  
→ Deadlock (neither can proceed)
```

Detection Algorithms [\[28\]](#) [\[29\]](#) [\[30\]](#) [\[27\]](#)

Three-Level Detection (Badal Algorithm, 1986): [\[28\]](#)

Level 1: Local deadlock check (same site)

- Fastest, no network I/O
- Detects 80% of deadlocks (cycles of length 2) [\[28\]](#)
- Check wait-for-graph (WFG) locally

Level 2: Remote resource check

- Combine lock histories from multiple sites
- Detect cycles involving 2-3 agents

Level 3: Global WFG construction

- Slowest, highest overhead
- Only used if Level 1/2 fail
- Guaranteed to detect all deadlocks [\[28\]](#)

Black-Box Monitors (2025): [\[29\]](#)

- Deploy monitors alongside each agent
- Monitors observe incoming/outgoing messages only
- Exchange "probes" to detect deadlocks
- **Proven:** Sound and complete (no false positives/negatives) [\[29\]](#)

Prevention Strategies [\[31\]](#) [\[30\]](#)

Timeout-Based (Recommended):

```
acquired = await redis.set(f"lock:{resource}", agent_id, ex=60, nx=True)
if not acquired:
    # Timeout after 60 seconds
    raise LockTimeout()
```

Priority-Based Preemption:

```
if current_holder_priority < self.priority:
    # Preempt lower-priority agent
    await redis.delete(f"lock:{resource}")
    acquired = await redis.set(f"lock:{resource}", self.id, ex=60, nx=True)
```

Resource Ordering:

```
# Always acquire locks in order: Lock1, Lock2, Lock3
locks = sorted(required_locks) # Alphabetical order
for lock in locks:
    await acquire_lock(lock)
```

Recommendation: Timeout-based locks (60s) + priority preemption (Orchestrator > Specialists).

Final Architecture Decision Matrix

Orchestration Pattern

| Pattern | Use Case | Complexity | Scalability | Verdict |
|-------------------------|------------------------------|------------|-------------|--------------------------------|
| Pure Orchestration | Critical workflow path | Low | Medium | ✓ For test generation pipeline |
| Pure Choreography | Side effects, notifications | High | High | ✓ For event bus |
| Hybrid (Recommended) | Both critical + side effects | Medium | High | *** Best for Phase 3 |

Workflow Engine Comparison

| Feature | Netflix Conductor | Temporal | LangGraph | Custom (Redis) |
|---------------------|------------------------|----------------------|---------------|----------------|
| DSL | JSON | Code (deterministic) | Code (Python) | Code (Python) |
| Infrastructure | Heavy (server + queue) | Heavy (cluster) | Light (Redis) | Light (Redis) |
| Learning Curve | Medium | High | Medium | Low |
| LLM Compatible | ✓ Yes | ✗ No (determinism) | ✓ Yes | ✓ Yes |
| Time Travel Testing | ✗ No | ✓ Yes | △ Partial | ✗ No |
| Production Scale | *** Netflix | *** Uber | ** Growing | ★ Custom |
| Recommendation | △ Overkill | ✗ Incompatible | *** Best | ** Fallback |

Task Allocation Evolution

| Phase | Strategy | Agents | Complexity | When |
|---------|-----------------------|--------------|------------|---------------|
| Phase 1 | Direct Assignment | 1 per type | Low | Sprints 7-8 |
| Phase 2 | Weighted Round Robin | 2-3 per type | Medium | Sprints 9-10 |
| Phase 3 | Contract Net Protocol | 3+ per type | High | Sprints 11-12 |

Key Takeaways

Critical Decisions

1. Hybrid Orchestration (Centralized + Decentralized) [\[1\]](#) [\[2\]](#)

- **Orchestrated:** Critical workflow (Observe → Analyze → Evolve → Report)
- **Choreographed:** Event notifications (status, metrics, alerts)
- **Why:** Balance predictability (orchestration) with scalability (choreography)

2. LangGraph Over Conductor/Temporal

- **Conductor:** Too heavy (requires Conductor server + Kafka/SQS)
- **Temporal:** Incompatible (determinism requirement vs LLM non-determinism)
- **LangGraph:** Right-sized (checkpointing, retry, state machines with minimal infra)

3. Phased Task Allocation

- **Start:** Direct assignment (simple, fast)
- **Scale:** Weighted round robin (load balancing)
- **Optimize:** Contract Net Protocol (quality + resource awareness)

4. Registry + Health Monitoring

- **Registry:** Redis-based Agent Cards with heartbeat (30s)
- **Health Checks:** Liveness + Readiness + Deep
- **Graceful Shutdown:** 5-step process (stop → complete → checkpoint → release → deregister)

5. Deadlock Prevention

- **Timeout-based locks:** 60s max hold time
- **Priority preemption:** Orchestrator > Specialists
- **Resource ordering:** Alphabetical lock acquisition

Production Readiness Checklist

Orchestration:

- ✓ Hybrid pattern (orchestrated critical path + choreographed events)
- ✓ State machine for test generation workflow
- ✓ Retry logic with exponential backoff
- ✓ Human-in-the-loop pause/resume

Task Allocation:

- ✓ Phase 1: Direct assignment (Sprints 7-8)

- ✓ Phase 2: Weighted round robin (Sprints 9-10)
- ✓ Phase 3: Contract Net Protocol (Sprints 11-12)

Agent Lifecycle:

- ✓ Self-registration on startup
- ✓ Heartbeat every 30 seconds
- ✓ Graceful shutdown (5-step process)
- ✓ Auto-recovery (Kubernetes restart)

Monitoring:

- ✓ Health checks (liveness, readiness, deep)
- ✓ Metrics (availability, latency, error rate, queue depth)
- ✓ Alerts (critical: agent down >5 min, error rate >10%)

Deadlock Prevention:

- ✓ Timeout-based locks (60s)
- ✓ Priority preemption
- ✓ Exponential backoff on contention

This hybrid orchestration architecture, validated by Netflix's production experience and academic research, provides the **optimal balance** of control, scalability, and operational simplicity for Phase 3 multi-agent test automation.

**

1. <https://camunda.com/blog/2023/02/orchestration-vs-choreography/>
2. <https://www.processmaker.com/blog/process-orchestration-vs-choreography-microservices/>
3. <https://bytebytogo.com/guides/orchestration-vs-choreography-microservices/>
4. <https://www.tatvasoft.com/outsourcing/2024/06/microservices-orchestration-vs-choreography.html>
5. <https://www.index.dev/blog/choreography-vs-orchestration-microservices>
6. <https://www.getmaxim.ai/articles/best-practices-for-building-production-ready-multi-agent-systems/>
7. <https://www.geeksforgeeks.org/system-design/netflix-conductor-microservices-orchestration/>
8. https://netflixtechblog.com/netflix-conductor-a-microservices-orchestrator-2e8d4771bf40?gi=74f811c1_62c9
9. <https://news.ycombinator.com/item?id=24214735>
10. <https://temporal.io/blog/workflow-engine-principles>
11. <https://www.baeldung.com/spring-boot-temporal-workflow-engine>
12. <https://thesai.org/Publications/ViewPaper?Volume=4&Issue=11&Code=IJACSA&SerialNo=6>
13. <https://www.sciencedirect.com/topics/computer-science/contract-net-protocol>
14. <https://arxiv.org/abs/2601.08815>
15. <https://microservices.io/patterns/service-registry.html>

16. <https://www.solo.io/topics/microservices/microservices-service-discovery>
17. <https://görn.name/posts/buildingagentdiscovery/>
18. <https://www.truefoundry.com/blog/ai-agent-registry>
19. <https://dev.to/vishnuchilamakuru/4-service-discovery-patterns-for-your-microservice-pg>
20. <https://www.geeksforgeeks.org/system-design/graceful-shutdown-in-distributed-systems-and-microservices/>
21. <https://www.linkedin.com/pulse/mastering-graceful-shutdown-distributed-systems-jainal-gosaliya-oliye>
22. <https://nelsonbn.com/blog/dotnet-graceful-shutdown/>
23. <https://docs.aws.amazon.com/streams/latest/dev/agent-health.html>
24. <https://mbrenndoerfer.com/writing/monitoring-reliability-ai-agents>
25. <https://docs.aws.amazon.com/firehose/latest/dev/agent-health.html>
26. <https://oasis.library.unlv.edu/cgi/viewcontent.cgi?article=1170&context=rtds>
27. <https://ics.uci.edu/~cs223/papers/p303-knapp.pdf>
28. <https://dl.acm.org/doi/pdf/10.1145/6513.6516>
29. <https://arxiv.org/html/2508.14851v3>
30. https://www.ijtra.com/view.php_paper-study-of-diverse-models-of-deadlock-detection-in-distributed-environment.pdf
31. <https://www.ifaamas.org/Proceedings/aamas2024/pdfs/p580.pdf>
32. [Phase3-Developer-A-Research-Plan.md](#)
33. <https://www.geeksforgeeks.org/system-design/orchestration-vs-choreography/>
34. <https://www.wallarm.com/what/orchestration-vs-choreography>
35. <https://www.di.univr.it/documenti/Occorrenzialns/matdid/matdid610023.pdf>
36. <https://www.milanjovanovic.tech/blog/orchestration-vs-choreography>
37. <https://stackoverflow.com/questions/51680291/netflix-conductor-as-a-workflow-engine-solution>
38. <https://developer.hashicorp.com/consul/docs/use-case/service-discovery>
39. <https://wseas.com/journals/articles.php?id=10940>
40. <https://fiveable.me/swarm-intelligence-and-robotics/unit-5/consensus-algorithms/study-guide/PWtgaBKcesTRA9TJ>
41. <https://www.reateai.com/blog/task-scheduling-algorithms-in-distributed-systems/84c0c02be9e5cc6114cf22ab73c7981c>
42. <https://arxiv.org/html/2205.12880v2>
43. <https://www.geeksforgeeks.org/java/service-discovery-and-service-registry-in-microservices/>
44. <https://www.cse.ust.hk/~psander/docs/alloc.pdf>
45. <https://arxiv.org/html/2504.02128v1>
46. <https://www.f5.com/company/blog/nginx/service-discovery-in-a-microservices-architecture>
47. <https://dl.acm.org/doi/10.1145/800174.809792>
48. <https://pmc.ncbi.nlm.nih.gov/articles/PMC9024934/>
49. <http://www.isca.me/rjrs/archive/v1/IISC-2011/74.ISCA-ISC-2011-5CITS-15.pdf>
50. <http://anale-informatica.tibiscus.ro/download/lucrari/8-2-02-Opeyemi.pdf>

